

SERAMAN - AI VIDEO PRODUCTION SYSTEM

System Handoff v2

Updated delivery documentation. Reflects all changes made since the initial handoff.


COMPLETE DELIVERY

DELIVERED	VERSION	PLATFORM	BUILT BY
July 2026	2.0	n8n + Kie AI + Creatomate + Biotato	Emmanuel Adekoya

WHAT YOU NOW HAVE

Submit Once. The System Does the Rest.

Drop a product image and a description into the form. Eight scene images are generated and sent to you for review. Flag anything that looks wrong and only those scenes are regenerated. Once every image is approved, eight video clips are produced. The final video is assembled, sent to you for review, and published to Instagram Reels, TikTok, YouTube Shorts, and Facebook the moment you approve. You are notified at every stage by email.



Your only inputs: the product image, the product description, and two approvals. Everything between them runs automatically.

HOW TO USE IT

What You Actually Do

Four steps. Two of them are clicking a button to approve what the system already built.

- 1 Open your product form**
This is your permanent submission link. Bookmark it. Every new product starts here.
<https://tally.so/r/r/bx5vx>
- 2 Upload the product image and description**
One product image is required. A second optional angle can be added for variety between scenes. The Italian product description goes in the text field. Submit once and the pipeline starts immediately.
- 3 Review your 8 scene images**
Within minutes you receive an email with all 8 generated images for that product. If any image looks wrong, flag those specific scenes on the linked form. The system regenerates only the flagged ones and sends updated images. If everything looks good, approve all. Video generation starts only after every image is approved.
- 4 Review and approve the final video**
Once all 8 video clips are assembled, you receive a final review email with a direct link to watch the video. If it is ready, approve on the form. The system publishes to all four platforms automatically and sends a delivery confirmation to your inbox.

SYSTEM ARCHITECTURE

Six Workflows, One System

The pipeline is split across six modular n8n workflows. Each handles a distinct stage of the process. This means any single stage can be updated, replaced, or scaled without touching the others.

WORKFLOW 01
SERAMAN Product Automation
The master orchestrator. Receives the Tally form submission, validates input, deduplicates repeat submissions, runs the Claude script agent, and calls each sub-workflow in sequence.


WORKFLOW 02
SERAMAN Generate Images
Scene image generation. Generates all 8 scene images via Kie AI, polls for completion with retry logic, writes results to the tracking sheet, and sends the image review email to your inbox.

WORKFLOW 03
SERAMAN Scene Approval
Approval gate and regen router. Receives your approve or flag response, routes flagged scenes for regeneration, sends the pipeline to video generation once all scenes are approved.

WORKFLOW 04
Seraman Generate Videos
Dual-branch parallel video generation engine. Runs Branch A (scenes 1 and 8) and Branch B (scenes 2 through 7) simultaneously, polls Kie AI for results with retry logic, and writes all video URLs to the tracking sheet.

WORKFLOW 05
SERAMAN Error Handler
Global error monitor. Fires automatically when any workflow in the n8n instance encounters an error. Sends a branded HTML email with the exact failure details and a direct link to the execution. Distinguishes platform outages from actual pipeline errors.

WORKFLOW 06
Seraman Post to Socials
Publishing and delivery. Receives the approved final video, generates platform-optimized captions for each channel, uploads to Biotato, publishes simultaneously to all four platforms, and sends a delivery confirmation email.




↑ SERAMAN Product Automation - master orchestration workflow

IMAGE GENERATION + APPROVAL GATE

Quality Control Before a Single Video Credit Is Spent

Image generation runs first, before any video is produced. All 8 scene images are sent to you for review. You flag anything that looks wrong — wrong product form, hallucinated text, incorrect scale — and only those scenes are regenerated. Video generation does not start until every image is approved.

This eliminates the most expensive failure mode in AI video pipelines: generating 8 full video clips from a bad reference image, then discarding them all. Catching the problem at the image stage costs a fraction of catching it after video generation.




↑ Generate Images - async image generation with per-scene polling and retry

↑ Scene Approval - approve or flag routing, regen loop, pipeline release gate

VIDEO GENERATION + ASSEMBLY

Generate, Assemble, Deliver

The Generate Videos workflow runs two branches in parallel — Branch A for scenes 1 and 8, Branch B for scenes 2 through 7 — each with its own async polling state machine. Once all 8 clips are confirmed complete, the Edit Videos workflow builds the Creatomate render payload, submits assemblies, and polls until the final MP4 is ready. The assembled video is sent to you for review before anything is posted publicly.



↑ Generate Videos - dual-branch Kie AI generation with self-healing retry

↑ Post to Socials - Biotato publishing to 4 platforms on approval

YOUR EMAIL SEQUENCE

Five Emails. One Per Stage.

You receive a branded email at every meaningful checkpoint. Nothing requires you to log into n8n to check status.


IMAGES_READY
All 8 scene images are generated. The email shows every image for review and includes a direct link to the approve or flag form. This fires before video generation starts.

IMAGES_UPDATED
Flagged scene images have been regenerated. The email shows only the updated images. Approve all to release the pipeline to video generation.

VIDEO_READY
The final assembled video is ready for review. Direct link to watch, with an approve button. Approving this triggers publishing to all four platforms.

VIDEO_APPROVED
The pipeline received your approval and is now publishing to all four platforms. No action needed.

POSTS_SCHEDULED
Publishing confirmed. Posts are live or scheduled across Instagram Reels, TikTok, YouTube Shorts, and Facebook. Delivery is complete.



↑ IMAGES_READY - scene image review email with approve or flag link

↑ VIDEO_READY - final video review email with direct watch and approve links

TRACKING & LOGGING

Your Google Sheet

Every pipeline run is logged across three sheets. You can check the status of any product run at any time without logging into n8n.

Sheet 1	Scene generation tracking. Every scene's image prompt, video prompt, reference image URL, product image URLs, generated image URL, and status. Updated in real time as each scene completes. One row per scene per product run.
Sheet 2	Scene video URLs. The raw Kie AI video output URL for each scene, keyed by scene number and job ID. Used by the Edit Videos workflow to build the Creatomate assembly payload.
Sheet 3	Final video log. The assembled Creatomate video URL, the original Italian product description, image regeneration round count, and final status. One row per completed product run.

ENGINEERING DEPTH

Under the Hood

What follows is a summary of the non-obvious engineering decisions built into this system — the problems that do not appear in demos but surface in production at scale.

HUMAN-IN-THE-LOOP QUALITY GATE - BEFORE ANY VIDEO CREDIT IS SPENT
The image approval gate is the most significant change in v2. AI image generation can produce wrong product forms, hallucinated text, or incorrect scale — problems that are cheap to catch on a static image but expensive to discover after full video generation. The gate intercepts the pipeline at the image stage: 8 images are generated and reviewed before a single video clip is submitted to Kie. Flagged scenes are regenerated individually. The pipeline only continues when every image is explicitly approved by you. No video credit is ever burned on a bad source image.

ASYNCHRONOUS STATE MACHINES WITH SELF-HEALING RETRY
Both image and video generation are asynchronous — submitting a scene returns a task ID, not a result. The system runs an independent polling state machine per scene that checks every 30 to 60 seconds. Each scene tracks two counters: poll attempts (up to 20, roughly 10 to 60 minutes of patience per scene) and regen attempts (up to 3 full resubmissions if Kie reports a generation failure). A three-way router classifies every poll response into success, failed, or pending and routes each path differently. Failed scenes are flagged with a reason and written to the sheet. They never disappear silently.

ITEM IDENTITY INTEGRITY ACROSS ASYNC LOOPS
n8n's default item-linking is positional — it assumes item N in maps to item N out. This breaks when a failed scene gets resubmitted and takes a different path through the workflow at a different time. The fix: scene_number travels forward as explicit data on every item through every node. Results are always attributed to the correct scene regardless of what order they return from Kie AI. This eliminates an entire class of silent misattribution bugs that would cause the wrong video clip to appear in the wrong scene slot.

CLAUDE V5.10 SCRIPT ENGINEERING
The Claude script agent is now on version 5.10. It enforces a Problem to Experience to Feature to Benefit formula per scene, uses scene chaining so each scene connects naturally to the previous, and applies buyer psychology ordering across the feature sequence. A full banned language list hard-blocks terms that trigger skepticism in tactical audiences. Version 5.10 adds automatic phonetic respelling for English brand and technical terms inside Italian voiceover — Cerakote becomes Cerakot, so Italian text-to-speech reads it correctly without any manual correction.
→ Strict JSON schema output — no markdown, no preamble, machine-parseable
→ 12-word hard cap per voiceover line with finish-by-second-six pacing beats
→ Scene chaining — every scene opens with a connector to the previous
→ Phonetic respelling — English loanwords automatically converted for Italian TTS

FAULT TOLERANCE - NO SCENE IS EVER LOST
If one scene out of eight fails irrecoverably at either the image or video stage, the remaining scenes still complete and are logged correctly. Both branches converge through a Merge node, get sorted back into correct scene order, and write to the output sheet as a complete set. Failed scenes are clearly marked for review. The entire batch never errors out because of a single scene failure.

ERROR HANDLING

Global Error Monitor

The Error Handler workflow runs independently of the main pipeline. It fires automatically whenever any workflow in the n8n instance throws an error. The email it sends is fully branded and built for debugging: JetBrains Mono font, dark matte design, highlighted error box, and a one-click link directly to the failed execution.

The error handler distinguishes between two failure types. A red alert means an actual pipeline error that needs investigation. An amber alert means the issue is with an upstream provider — Kie AI, Creatomate, or Biotato — and the pipeline itself is healthy. You will never wonder whether your automation is broken or whether a platform had an outage.

↑ SERAMAN Error Handler - global monitor for all workflow failures

YOUR FORMS

Two Tally Forms. Both Connected to the Pipeline.

The pipeline uses two Tally forms. One to submit products. One to approve or flag scenes. Responses trigger the next stage automatically.

PRODUCT SUBMISSION - TALLY.SO/R/0BX5VX
Your permanent product form. Upload the product image, an optional second angle, and the Italian product description. Submit once to trigger the full pipeline. Bookmark this link.

SCENE APPROVAL - UNIQUE LINK PER JOB
Sent in each image and video review email. Checkboxes for scenes 1 through 8. Check any scene to flag it for regeneration. Check Approve All to release the pipeline. Do not bookmark this link — it changes per job.

↑ Product submission form - permanent link, bookmark this

↑ Scene approval form - unique per job, linked in every review email

PUBLISHING

Platform Publishing via Biotato

On approval of the final video, the assembled clip and Italian captions are passed to Biotato, which publishes simultaneously to all four platforms in a single operation. Each platform receives a caption engineered for its algorithm.
Instagram Reels: vertical 9:16, caption with 125-character hook before the fold, hashtags positioned for the Reels algorithm. TikTok: first line under 100 characters, hook matched to the product type, three to five niche tags only. YouTube Shorts: title optimized for search, 60 characters max, product name plus key differentiator. Facebook: longer copy with an engagement question, optimized for comments over views.

RUNNING COSTS

Monthly Platform Costs

These costs are paid directly by Seraman. They are not included in the build fee.

SERVICE	COST
Kie AI — image generation + video generation (Veo 3.1)	Per-sec-use • ~\$3-8 per finished video
Creatomate — video assembly	~\$29/mo
Biotato — social publishing, all 4 platforms	\$29/mo flat • no per-post fees
n8n Cloud — workflow orchestration	~\$28/mo • or self-hosted free
Tally — product intake and scene approval forms	Free tier sufficient
Total platform cost (excluding generation)	~\$78-88/mo

At 10 videos per month, total cost including generation stays under \$130/mo. Generation cost scales directly with volume — no fixed commitment.

A model upgrade to Gemini Omni has been tested and is available. It resolves a specific class of joint and hand defects in motion scenes and produces more natural voiceover than Veo 3.1 Fast. The per-clip cost difference is approximately \$1.50 to \$2.00 per finished video at 8 scenes. The recommendation to upgrade stands when you are ready.

YOUR ACCESS

What You Own and Control

All platform accounts are registered in Seraman's name. Below is a reference for what you have access to and where each service is managed.

Product Submission Form	tally.so/r/r/bx5vx — bookmark this link
Scene Approval Form	Unique link per job — sent in each review email, do not bookmark
Workflow Platform	n8n — all six workflows visible under Personal projects
Tracking Sheet	Google Sheets — three tabs: scene generation, scene video URLs, final videos
Image and Video Generation	Kie AI — credits loaded as needed, no monthly subscription
Video Assembly	Creatomate — template ID stored in the Edit Videos workflow
Social Publishing	Biotato — all four platform accounts connected

WHAT'S NEXT

Long-Form Content Module

The core pipeline is now in production. Image generation, approval gate, video generation, assembly, and publishing — all automated. The natural next build is long-form content: 5 to 8 minute videos for YouTube and deep-funnel campaigns. These convert differently from short Reels — they build purchase confidence before someone clicks. The current pipeline architecture already supports this layer. When you are ready to build it, nothing needs to start from scratch.

Giovanni — la versione 2 di questo sistema è più robusta, più intelligente e più umana di tutto quello che avevamo immaginato quando abbiamo iniziato. È stato un piacere costruirlo. In bocca al lupo per i prossimi lanci.